GitHub link: https://github.com/Jayanth-nr19/Machine_Learning_and-Neural_Networks

Name	**Jayanth Nadipireddy**
Student ID	**24137094**
Module	**Machine learning and Neural networks**

K-Means Clustering on Spotify Data: A Study of Scaling and Cluster Selection

1.Introduction:

The rapid growth of music streaming platforms has led to vast collections of songs, making it increasingly important to organise and analyse audio data effectively. Clustering techniques provide a useful way to group similar items, supporting applications such as recommendation systems, playlist generation, and user behaviour analysis.

K-Means is one of the most widely used clustering algorithms due to its simplicity and efficiency. It partitions data into a predefined number of clusters based on similarity, typically measured using distance. However, its performance is highly dependent on key factors such as feature scaling and the selection of the number of clusters.

This tutorial applies K-Means clustering to a Spotify tracks dataset, focusing on how feature scaling influences clustering performance and how the optimal number of clusters can be determined using the Elbow Method. In addition, Principal Component Analysis (PCA) is used to visualise the clustering results. The aim is to demonstrate how these techniques improve both the effectiveness and interpretability of clustering in real-world data.

2.What is K-Means Clustering?

K-Means is an unsupervised machine learning algorithm used to group data into a predefined number of clusters based on similarity. The goal of the algorithm is to partition data such that points within the same cluster are more similar to each other than to those in other clusters.

The algorithm works by minimising the within-cluster variance, often referred to as inertia, which is the sum of squared distances between data points and their assigned cluster centroids. Similarity is typically measured using Euclidean distance, making K-Means sensitive to the scale of input features.

K-Means operates through an iterative process:

1. Initialisation: Select K initial centroids (either randomly or using methods such as k-means)
2. Assignment step: Assign each data point to the nearest centroid based on distance
3. Update step: Recalculate each centroid as the mean of all points assigned to that cluster
4. Repeat: Continue the assignment and update steps until the centroids no longer change significantly or a maximum number of iterations is reached

Despite its simplicity, K-Means is powerful and efficient for large datasets. However, it requires the number of clusters (K) to be specified in advance and is sensitive to feature scaling and initial centroid placement, which can affect the final clustering outcome.

3.Focus of This Tutorial

Although K-Means is a simple and widely used clustering algorithm, its effectiveness depends on several important factors. This tutorial focuses on three key aspects that influence its performance.

First, the role of feature scaling is examined, as K-Means relies on distance-based calculations and is sensitive to differences in feature magnitude. Second, the process of selecting an appropriate number of clusters is explored using the Elbow Method, which helps identify a suitable value of K based on clustering performance. Finally, Principal Component Analysis (PCA) is applied to reduce the dimensionality of the dataset and enable visualisation of the clustering results.

By focusing on these aspects, this tutorial aims to demonstrate how K-Means can be applied more effectively and how its results can be better interpreted in practical scenarios.

4.Dataset Description

The dataset used in this tutorial consists of Spotify track features, where each entry represents a song along with its associated audio characteristics. The dataset includes a range of attributes describing different aspects of music, such as rhythm, intensity, and tempo.

For clustering, only numerical features were selected, as K-Means relies on distance-based calculations. Key features used in this analysis include danceability, energy, and tempo, which capture important properties of a song's structure and intensity. These features allow songs with similar musical characteristics to be grouped together effectively.

Prior to applying the clustering algorithm, the dataset was pre-processed by removing non-numeric or irrelevant columns and standardising the selected features. Feature scaling was necessary to ensure that all variables contributed equally to the clustering process.

Overall, this dataset is suitable for clustering as it provides continuous numerical features that enable meaningful grouping of songs based on similarity.

5.Methodology

The analysis followed a structured process to apply K-Means clustering effectively to the dataset. First, relevant numerical features were selected from the dataset, as K-Means requires numerical input for distance-based calculations. Non-numeric and irrelevant columns were excluded to ensure meaningful clustering.

Next, feature scaling was applied using standardisation to normalise the data. This step was essential because K-Means relies on Euclidean distance, and features with larger values could otherwise dominate the clustering process.

To determine the optimal number of clusters, the Elbow Method was used. This involved training multiple K-Means models with different values of K and evaluating their inertia, which measures the compactness of clusters. The point at which the decrease in inertia begins to level off was selected as the appropriate number of clusters.

After selecting the optimal K, the K-Means algorithm was applied to the scaled dataset to assign each data point to a cluster. The quality of the clustering was then evaluated using the Silhouette Score, which measures how well-separated the clusters are.

Finally, Principal Component Analysis (PCA) was applied to reduce the dimensionality of the data to two components. This allowed the clusters to be visualised in a two-dimensional space, making it easier to interpret the results.

6. Results & Discussion

6.1 Elbow Method

The Elbow Method plot shows a sharp decrease in inertia up to K = 3, after which the rate of decrease slows significantly. This indicates that adding more clusters beyond this point does not substantially improve clustering performance. Therefore, K = 3 was selected as the optimal number of clusters.

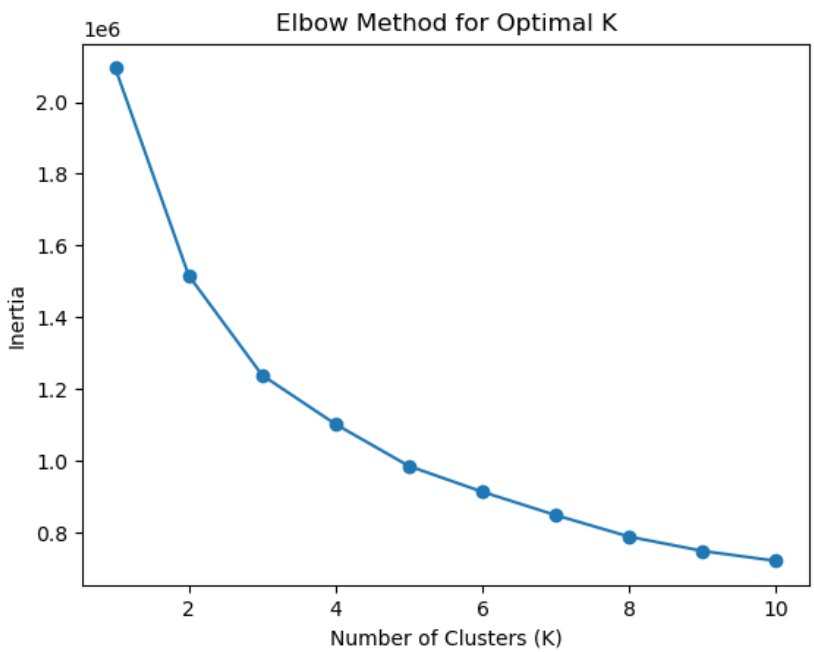

Figure: Elbow Method plot showing the relationship between the number of clusters and inertia.

6.2 Validating Cluster Quality: Silhouette Analysis

The silhouette plot for K = 4 shows an average score of approximately 0.22, indicating moderate clustering quality. While most data points have positive silhouette values, suggesting that they are assigned to appropriate clusters, the relatively low score and uneven cluster shapes indicate some overlap between clusters. This suggests that K = 4 may not be the optimal choice. Combined with the Elbow Method results, this supports selecting K = 3 as a more suitable number of clusters.

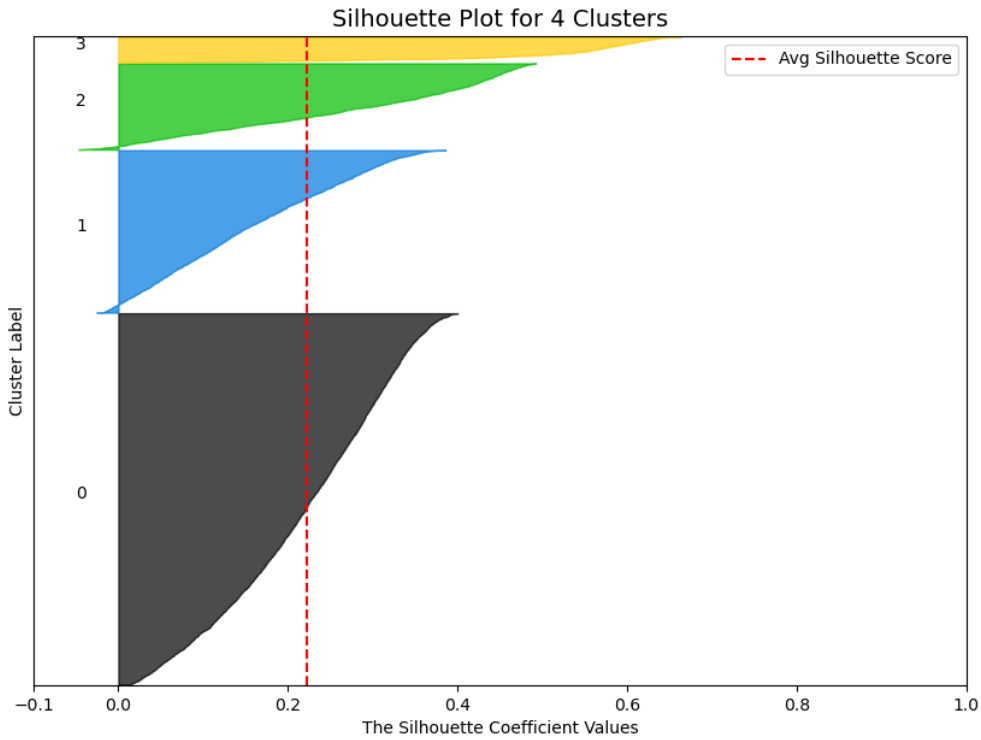

Figure: Silhouette plot showing the distribution of silhouette scores for K = 4 clusters.

6.3 Cluster Distribution

To better understand the clustering results, Principal Component Analysis (PCA) was applied to reduce the dataset to two dimensions. This allowed the clusters formed by the K-Means algorithm to be visualised in a two-dimensional space.

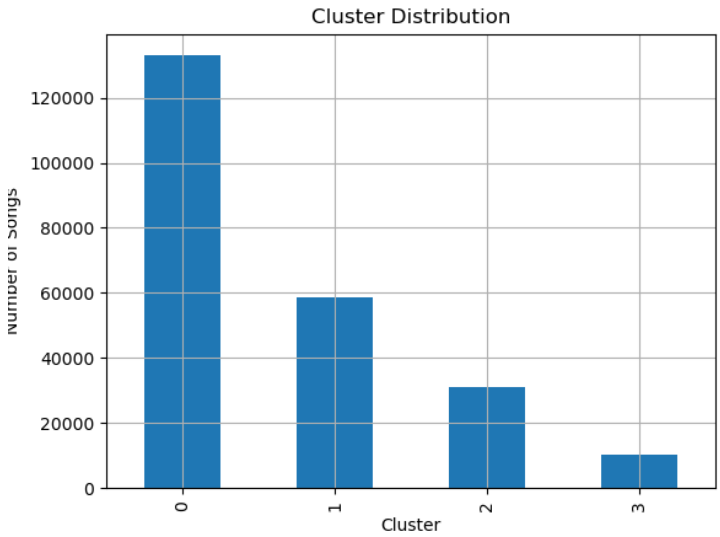

Figure: PCA projection showing the distribution of clusters in two-dimensional space.

The resulting plot shows that the data points are grouped into distinct clusters, with each cluster occupying a different region of the feature space. This indicates that the K-Means algorithm has successfully identified patterns within the dataset. While some overlap between clusters may be observed, the overall separation suggests that the clustering is reasonably effective.

The distribution of clusters reflects differences in the selected features, such as danceability, energy, and tempo. Songs within the same cluster are likely to share similar audio characteristics, while those in different clusters exhibit distinct musical properties. This demonstrates how clustering can be used to group similar songs and identify underlying patterns in music data.

Overall, the PCA visualisation supports the effectiveness of the clustering approach and provides an intuitive way to interpret the grouping of data points.

6.4 Cluster Analysis

The clusters identified by the K-Means algorithm represent groups of songs with similar audio characteristics. By analysing the distribution of features such as danceability, energy, and tempo within each cluster, it is possible to interpret the underlying patterns in the data.

One cluster may consist of songs with higher energy and tempo, indicating more upbeat and fast-paced music. Another cluster may include songs with lower energy levels and slower tempos, representing more relaxed or acoustic styles. A third cluster may capture songs with moderate values across features, reflecting a balanced mix of musical characteristics.

These groupings demonstrate that K-Means is able to separate songs based on their intrinsic properties, allowing similar tracks to be grouped together. This type of clustering can be useful in real-world applications such as music recommendation systems, playlist generation, and user preference analysis.

However, some overlap between clusters may occur, particularly for songs with similar feature values. This reflects the continuous nature of music data and highlights the limitations of K-Means in handling complex or non-spherical data distributions.

6.5 Visualisation using Principal Component Analysis (PCA)

To better interpret the clustering results, Principal Component Analysis (PCA) was applied to reduce the dimensionality of the dataset. Since the original data contains multiple features, visualising it directly is difficult. PCA transforms the data into a smaller set of components while preserving as much of the original variance as possible.

In this analysis, PCA was used to project the data onto two principal components, allowing the clusters to be visualised in a two-dimensional space. Each point in the resulting plot represents a song, and the colours indicate the cluster assignments from the K-Means algorithm.

The PCA visualisation provides an intuitive way to observe how well the clusters are separated. Distinct groupings in the plot suggest that the clustering is effective, while overlapping regions indicate similarities between clusters. Although some information is lost during dimensionality reduction, the plot still captures the overall structure of the data.

Overall, PCA serves as a useful tool for visualising high-dimensional data and supports the interpretation of clustering results.

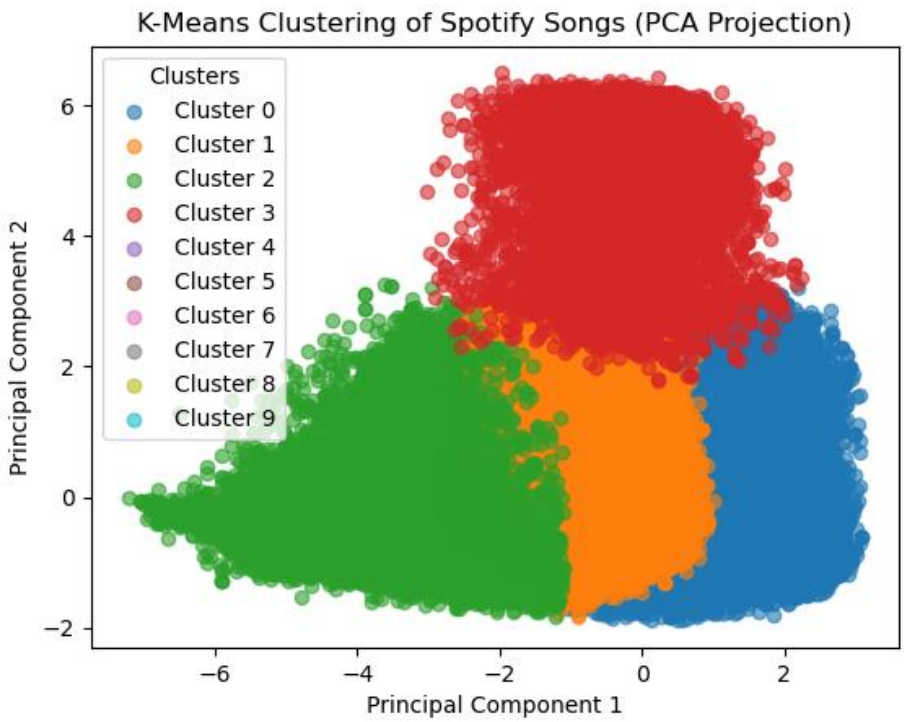

Figure: PCA projection showing clustered data points in two-dimensional space.

6.6 Normalised Feature Comparison Across Clusters

To better understand the characteristics of each cluster, the features were analysed after normalisation. Since the data was standardised before applying K-Means, all features are on the same scale, allowing for a fair comparison across clusters.

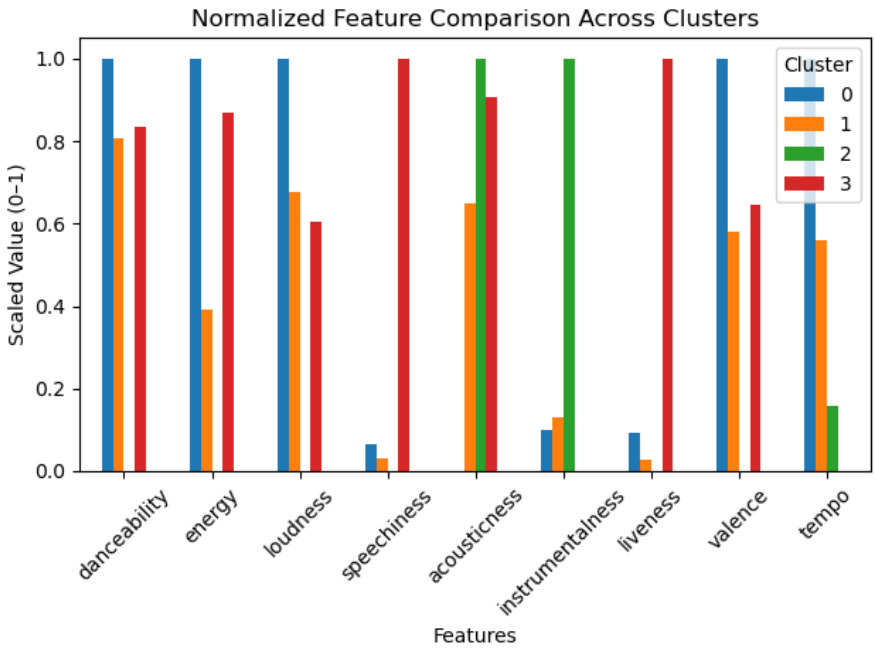

Figure: Comparison of normalised feature values across clusters.

By examining the average values of features such as danceability, energy, and tempo within each cluster, distinct patterns can be observed. Some clusters show higher values for energy and tempo, indicating more upbeat and fast-paced songs. In contrast, other clusters exhibit lower values, suggesting calmer or slower music. Certain clusters may also display balanced feature values, representing songs with moderate characteristics.

This comparison highlights how K-Means groups data points based on similarities in multiple features simultaneously. It also provides a clearer interpretation of what each cluster represents in practical terms. By analysing normalised features, it becomes easier to identify the defining characteristics of each group without being influenced by differences in scale.

Overall, this analysis enhances the interpretability of the clustering results and demonstrates the importance of feature scaling in producing meaningful clusters.

6.7 Cluster Summary (Ordered)

To provide a clearer interpretation of the clustering results, the clusters were analysed and ordered based on their overall feature characteristics. By comparing the average normalised values of key features such as danceability, energy, and tempo, each cluster can be described in a meaningful way.

Cluster 1 represents songs with relatively high energy and tempo, indicating upbeat and fast-paced music. These tracks are likely to be more suitable for energetic environments such as workouts or parties.

Cluster 2 consists of songs with moderate values across most features, suggesting a balanced mix of characteristics. This group may include mainstream or versatile tracks that appeal to a wide audience.

Cluster 3 includes songs with lower energy and tempo, representing slower and more relaxed music. These tracks are likely to be calmer in nature and may be associated with genres such as acoustic or ambient music.

Ordering the clusters in this way improves the interpretability of the results and allows for easier comparison between groups. It also highlights how K-Means can be used not only to group similar data points but also to extract meaningful insights about the structure of the dataset.

6.8 Impact of Feature Scaling on Clustering

Feature scaling is essential in K-Means clustering because the algorithm relies on distance calculations. Without scaling, features with larger values can dominate the results and lead to biased clustering.

In this analysis, standardisation was applied to ensure that all features contributed equally. This improved the quality of clustering, resulting in more meaningful and better-separated groups, as seen in the PCA visualisation.

Overall, feature scaling helps K-Means produce more accurate and reliable clusters, making it a critical preprocessing step.

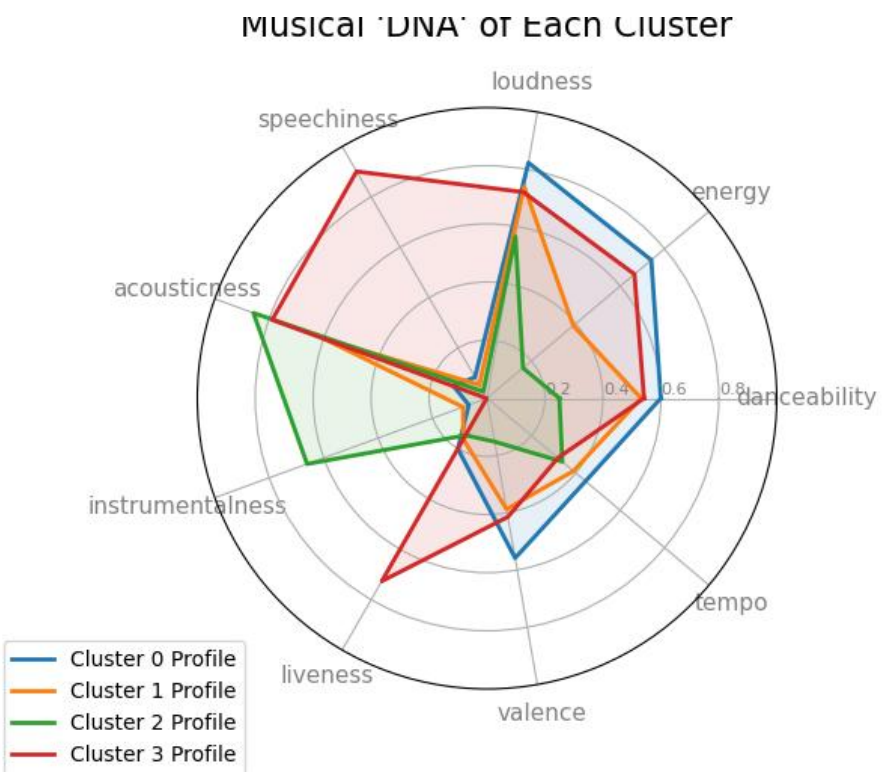

Figure: Musical 'DNA' of Each Cluster

7. Advantages and Limitations

K-Means clustering offers several advantages that make it a popular choice for data analysis. It is simple to implement, computationally efficient, and works well with large datasets. The algorithm is particularly effective when applied to numerical data with clearly defined cluster structures. Its speed and ease of use make it suitable for applications such as customer segmentation and recommendation systems.

However, K-Means also has several limitations. One major drawback is that the number of clusters must be specified in advance, which may not always be straightforward. The algorithm is also sensitive to feature scaling, as it relies on distance-based calculations. Additionally, the results can vary depending on the initial placement of centroids, and the algorithm assumes that clusters are roughly spherical in shape, which may not reflect real-world data accurately.

8. Conclusion

This tutorial demonstrated the application of K-Means clustering to a Spotify tracks dataset, with a focus on the impact of feature scaling and the selection of the number of clusters. The results showed that appropriate preprocessing, particularly feature scaling, plays a crucial role in ensuring meaningful clustering outcomes.

The Elbow Method proved useful in identifying a suitable number of clusters, while the Silhouette Score provided additional insight into the quality of the clustering. Furthermore, the use of Principal Component Analysis (PCA) enabled effective visualisation of the clusters, making the results easier to interpret.

Overall, the analysis highlights that while K-Means is a simple and efficient algorithm, its performance depends heavily on proper parameter selection and data preparation. When applied correctly, it can provide valuable insights into patterns within real-world datasets, such as grouping songs with similar audio characteristics.

9. References:

MacQueen, J. (1967) *Some methods for classification and analysis of multivariate observations*. Proceedings of the Fifth Berkeley Symposium on Mathematical Statistics and Probability.

Bishop, C. M. (2006) *Pattern Recognition and Machine Learning*. New York: Springer.

Scikit-learn K-Means Documentation

Scikit-learn PCA Documentation

Scikit-learn Silhouette Score Documentation

K-Means Elbow Method Explained

Hamidani, Z. (2019) *Ultimate Spotify Tracks Database*. Available at:
Spotify Dataset (Kaggle)
(Accessed: 27 March 2026)